


Low-Rank Few-Shot Adaptation of Vision–Language Models

Understanding, reproducing, and extending **CLIP-LoRA** — adapting CLIP from just a handful of labelled images.

CLIP · Radford 2021

CLIP-LoRA · Zarella 2024

paper: 11 datasets · ours: 10/11 reproduced



What we'll walk through

01 What is CLIP?

The vision–language model we adapt — how it's trained and how zero-shot classification works.

02 What is CLIP-LoRA?

Low-rank adapters inside CLIP's attention blocks — the design choices and why they work.

03 Our reproduction

Setup, scope, final results, and the challenges we hit.

04 Extension & outlook

A regularization term that anchors the adapted model to zero-shot CLIP — motivation, experiments, and results.

Adapting a giant model with only a few examples

Vision–language models like **CLIP** are pre-trained once on hundreds of millions of image–text pairs. To use one on **your** task you only have a handful of labelled images — the **few-shot** regime.

The question this paper answers: **what is the cheapest, most effective way to specialise CLIP from 1–16 examples per class** — without fine-tuning the whole network or running slow prompt-optimization?

FULL FINE-TUNING

Too many parameters

Millions of weights, easy to overfit on a few shots, heavy to store per task.

PROMPT TUNING

Slow to train

Methods like PLOT++ need many hours of optimization per task.

CLIP-LoRA

Few params • fast • accurate

Train tiny low-rank matrices inside attention — minutes, not hours.


What is CLIP?

The foundation model we're adapting — how it learns to connect images and language.




One model, two encoders, a shared space

- **Two encoders.** An **image encoder** (a Vision Transformer) and a **text encoder** (a Transformer) map their inputs into one shared embedding space.
- **Trained on scale.** ~400M image-text pairs scraped from the web — no manual class labels needed.
- **Aligned by meaning.** An image and its caption land **close together**; unrelated pairs land far apart.
- **The payoff.** Comparing image and text embeddings lets CLIP classify images it was never explicitly trained on — **zero-shot**.




How CLIP works, end to end


(1) Contrastive pre-training



(2) Create dataset classifier from label text



(3) Use for zero-shot prediction



1 Contrastive pre-training. Image & text encoders learn to match the correct pairs in each batch.

2 Build a classifier from labels. Class names go into a prompt → text embeddings.


3 Zero-shot prediction. Pick the class whose text embedding is most similar to the image.

Match the right caption out of the whole batch

For a batch of **N** image–text pairs, CLIP computes all **N × N** cosine similarities. The training objective pulls each image toward **its own** caption and pushes it away from the other N–1.

$$\text{sim}(I, T) = (I \cdot T) / (\|I\| \|T\|)$$

A symmetric cross-entropy loss over rows **and** columns maximises the **diagonal** (correct pairs) of the similarity matrix.



Classify by comparing to text prompts



PROMPT TEMPLATE – NO TRAINING, JUST TEXT

"a photo of a **cat** "

"a photo of a **dog** "

"a photo of a **forest** "

...one per class

Strong out of the box — but on **fine-grained** tasks (aircraft types, satellite terrain) zero-shot alone isn't enough. That's where few-shot **adaptation** comes in.

What is CLIP-LoRA, and how does it work?

From low-rank adapters to a fast, fixed recipe for few-shot adaptation.



Two families of few-shot adaptation

FAMILY A • PROMPT LEARNING

Optimise the text side

Learn continuous prompt tokens fed to the text encoder — **CoOp**, **CoCoOp**, **PLOT++**, **KgCoOp**, **ProGrad**, **MaPLe**.

Cost: back-propagation through the text encoder $\Rightarrow$ slow (hours).

FAMILY B • ADAPTERS

Add small trainable modules

Bolt lightweight layers onto frozen features — **CLIP-Adapter**, **Tip-Adapter-F**, **TaskRes**.

Limit: often touch only one modality; method-specific tuning.

CLIP-LoRA A third path. Inject **low-rank adapters directly into the attention weights** of **both** encoders. One fixed recipe, no per-task tuning, trains in minutes.


LoRA – freeze the big matrix, train a tiny one

Instead of updating a weight matrix **W**, LoRA **freezes it** and learns a **low-rank update** $\Delta W = \mathbf{B} \mathbf{A}$, where B and A are skinny matrices of rank $r \ll d$.

$$h = W_0 x + \mathbf{B} \mathbf{A} x$$

- **W₀ frozen** ❄️ — pretrained knowledge is preserved.
- **Only B, A train** 🔥 — a few thousand parameters per layer.
- **Zero inference cost** — B A can be merged back into W after training.

LOW-RANK DECOMPOSITION



rank $r = 2$ in CLIP-LoRA $\Rightarrow$ <1% of params trained

LoRA inside every attention block of both encoders

WHERE THE ADAPTERS GO

VISION ENCODER

Layer L	+LoRA
...	+LoRA
Layer 2	+LoRA
Layer 1	+LoRA

TEXT ENCODER

Layer L	+LoRA
...	+LoRA
Layer 2	+LoRA
Layer 1	+LoRA

every layer of **both** encoders carries adapters

Inside each self-attention block, the query / key / value projections are the levers. CLIP-LoRA attaches a low-rank update to **W_q, W_k, W_v**:

+ LoRA → W_q + LoRA → W_k + LoRA → W_v

- **Both modalities adapt.** Vision *and* text encoders move together — key for fine-grained tasks.
- **Every level.** Adapters sit at all encoder depths (the “All” placement).
- **One recipe, all tasks.** The same hyper-parameters are used for every dataset — no per-task search.

One fixed configuration, kept across the benchmark

HYPER-PARAMETER	VALUE
Rank r	2
Target matrices	W_q, W_k, W_v
Encoders	vision + text
Placement	all layers
Dropout p	0.25
Learning rate	$2e-4$
Scheduler	cosine
Batch size	32
Iterations	$500 \times \text{shots}$
Prompt	"a photo of a [class]"

The strength of the method is its **simplicity**: no per-dataset hyper-parameter search, no validation-set tricks. The same settings run across all reproduced datasets.

$r = 2$

tiny rank

24 GB

fits one GPU

Results are averaged over **3 random seeds (1, 2, 3)** — the convention we follow exactly in our reproduction.

Faster than prompt learning, competitive accuracy

TRAINING TIME • 16-SHOT IMAGENET • SINGLE A100

METHOD	TRAIN TIME
PLOT++	15 h 30
ProGrad	3 h 20
CoOp	2 h 00
CLIP-LoRA	50 min

- **≈ 18× faster** than PLOT++ to train, on the same task.
- **State-of-the-art accuracy** across the 11-dataset few-shot benchmark, with no per-task tuning.
- **Consistent across CLIP visual backbones** — ViT-B/16, ViT-B/32, ViT-L/14.
- **Our reproduction** (16-shot ImageNet, V100 / A40): ≈104 min B/16, 92 min B/32, 191 min L/14 — same method, slower GPU than the paper's A100. Whole project ≈ 275 GPU-h.

Takeaway: a generic, well-understood NLP/LLM technique — LoRA — transfers cleanly to vision–language adaptation and beats specialised methods on speed.

Reproducing the paper

Our setup, what we reproduced, the results — and the challenges along the way.



Built on the authors' code, our own runners

BASE

Official repo

Forked MaxZane11a/CLIP-LoRA and reproduced the default configuration faithfully.

PROTOCOL

Seeds 1·2·3

Every experiment averaged over three seeds, exactly as the paper reports.

HARDWARE

Single GPU

Single-GPU jobs run across both NVIDIA V100 and A40 (48 GB); every configuration fits on a single GPU.

OUR REPRODUCIBILITY SCAFFOLDING

run_sanity_check.sh

run_table_3.sh

run_table_4.sh

run_table_5.sh

run_fig3.sh

run_kl_ablation.sh

run_kl_table3.sh

aggregate_results.py

verify_results.py

Every run logs dataset, backbone, shots, seed, rank, targets, encoder, placement, dropout, accuracy, status into a single clip_lora_results.csv. KL runs use a separate clip_lora_kl.csv.

What we set out to reproduce

Main tables · ViT-B/16, B/32 & L/14

phase 1

CLIP-LoRA across **10 datasets** × shots {1,2,4,8,16} × 3 seeds.

$10 \times 5 \times 3 = 150$ runs per backbone (3 backbones)

Work split

parallel

Two of us run backbones in parallel; Fig.3 jobs split round-robin so rank/dataset load is balanced.

Figure 3 ablations

phase 2

Rank × attention-matrix set × encoder choice, plus placement — on ImageNet, StanfordCars, EuroSAT at 4-shot.

$129 \text{ combos} \times 3 \text{ datasets} = 387 \text{ runs}$ (1 seed)

Scope note

Full grid for **10 of the paper's 11 datasets**; SUN397 dropped because the original URL is dead and the available HuggingFace version is parquet-based, not the raw JPEG folder structure expected by the original loaders.

How we split the work

HB Hasan Berke Bankoglu (12432802)

- **Table 3** — ViT-B/16 main results
- **Table 5** — ViT-L/14 main results
- **Figure 3** — ImageNet & EuroSAT ablations
- **Extension** — KL regularization to zero-shot CLIP

ZA Zarek Asif (12432949)

- **Table 4** — ViT-B/32 main results
- **Figure 3** — StanfordCars ablations

Each backbone is presented on its own slide next — every reproduced average lands within $\approx 0.3\%$ of the paper across all shot counts and 10 datasets (SUN397 excluded).

Few-shot results · ViT-B/16

■ **ours** ■ **paper** Δ = ours-paper

Shots	ImgNet	Airc	EuroSAT	Cars	Food	Pets	Flowers	Caltech	DTD	UCF	Avg
0 zero-shot	66.7	23.8	42.1	65.6	85.3	88.2	67.4	93.3	44.5	65.2	64.2
1	70.4 70.4	29.8 30.2	71.1 72.3	69.8 70.1	84.7 84.3	91.9 92.3	82.3 83.2	93.9 93.7	54.0 54.3	76.0 76.3	72.4 72.7 -0.3
2	70.8 70.8	32.6 33.2	82.0 82.7	73.4 73.2	83.3 83.2	90.7 91.3	89.8 89.8	94.8 94.6	59.9 59.9	80.1 80.0	75.7 75.9 -0.1
4	71.5 71.4	38.0 37.9	85.3 84.9	77.1 77.4	83.0 82.7	90.5 91.0	93.3 93.7	95.0 95.2	63.7 63.8	81.1 81.1	77.8 77.9 -0.1
8	72.4 72.3	45.4 45.7	89.1 89.7	82.1 82.1	83.1 83.1	91.3 91.7	96.1 96.3	95.7 95.6	67.5 67.5	84.0 84.1	80.7 80.8 -0.1
16	73.6 73.6	54.9 54.7	92.3 92.1	86.3 86.3	84.1 84.2	92.3 92.4	97.9 98.0	96.2 96.4	71.8 72.0	86.4 86.7	83.6 83.6 0.0

Top-1 accuracy (%), 3-seed avg (seeds 1·2·3) · paper = Zanella & Ben Ayed (2024), Table 3 · averages over the 10 datasets we ran (SUN397 excluded) · row 0 = zero-shot CLIP (frozen, no adaptation)

Few-shot results · ViT-B/32

ours paper Δ = ours-paper

Shots	ImgNet	Airc	EuroSAT	Cars	Food	Pets	Flowers	Caltech	DTD	UCF	Avg
0 zero-shot	62.1	18.4	35.7	60.2	78.8	85.1	64.0	91.5	42.9	61.6	60.0
1	65.3 65.2	22.9 22.9	67.9 67.1	63.6 63.5	78.5 78.4	88.0 88.7	77.5 77.4	93.0 93.0	49.5 49.8	72.2 72.3	67.8 67.8 0.0
2	65.7 65.7	23.6 24.1	80.7 80.8	64.6 64.8	76.5 76.3	86.8 86.6	84.2 84.7	93.9 93.7	57.5 57.4	75.4 75.4	70.9 71.0 0.0
4	66.5 66.5	28.3 27.7	84.4 85.6	68.5 68.3	76.0 75.6	86.3 86.3	90.2 90.1	94.2 94.3	60.8 60.3	76.0 76.5	73.1 73.1 0.0
8	67.2 67.2	35.7 36.1	89.1 88.8	74.6 74.4	76.8 76.7	88.2 87.7	92.7 92.4	94.4 94.8	63.7 63.7	80.2 80.1	76.2 76.2 +0.1
16	68.4 68.4	46.1 44.9	91.9 91.8	79.9 79.7	78.1 78.2	89.0 88.8	96.0 96.2	95.3 95.2	69.0 68.2	82.2 82.8	79.6 79.4 +0.2

Top-1 accuracy (%), 3-seed avg (seeds 1-2-3) · paper = Zanella & Ben Ayed (2024), Table 4 · averages over the 10 datasets we ran (SUN397 excluded) · row 0 = zero-shot CLIP (frozen, no adaptation)




Few-shot results · ViT-L/14

■ **ours** ■ **paper** Δ = ours-paper

Shots	ImgNet	Airc	EuroSAT	Cars	Food	Pets	Flowers	Caltech	DTD	UCF	Avg
0 zero-shot	73.5	29.9	55.1	76.9	90.3	93.1	76.2	95.1	52.3	73.8	71.6
1	77.0 76.7	41.5 41.2	72.0 73.7	81.2 80.7	90.3 90.5	94.3 94.2	90.8 91.2	95.7 95.7	61.0 61.2	83.0 82.4	78.7 78.8 -0.1
2	77.4 77.3	43.2 42.4	85.7 85.9	82.7 82.3	89.9 89.8	93.8 93.2	95.1 94.7	96.9 96.8	67.4 67.3	84.8 84.2	81.7 81.4 +0.3
4	78.1 77.9	50.0 48.9	85.4 86.4	85.3 85.2	89.8 89.6	94.1 93.9	97.3 97.4	97.1 97.2	70.2 70.4	86.6 86.4	83.4 83.3 +0.1
8	78.7 78.5	57.9 57.5	89.5 90.0	88.8 88.7	89.5 89.7	94.1 94.2	98.3 98.0	97.0 97.0	72.7 72.2	88.7 88.3	85.5 85.4 +0.1
16	79.7 79.6	66.3 66.2	93.0 93.1	91.0 90.9	90.0 89.9	94.7 94.3	99.1 99.0	97.2 97.3	76.8 76.5	90.3 89.9	87.8 87.7 +0.2

Top-1 accuracy (%), 3-seed avg (seeds 1·2·3) · paper = Zanella & Ben Ayed (2024), Table 5 · averages over the 10 datasets we ran (SUN397 excluded) · row 0 = zero-shot CLIP (frozen, no adaptation)

Rank × matrices × encoder – reproduced



4-shot · top-1 accuracy (%) · one-seed reproduction of the paper's Figure 3 · flat across rank ⇒ default r=2 · QKV ≈ best matrix set · both encoders & all layers win

The ablation supports the four default choices

RANK

Saturates early

ImageNet QKV: **71.3** (r1) → 71.5 (r4) → **70.8** (r32) — flat, even dips at high rank (mild overfit).
⇒ **default rank = 2**

ATTENTION MATRICES

QKV is the sweet spot

rank2, both: Q 71.0 · V 71.4 · QV 71.6 · QKV 71.5 · QKVO 71.5 — adding O gives nothing. ⇒
 W_q, W_k, W_v

ENCODER

Adapt both

ImageNet QKV: vision 70.0 · text 71.0 · **both 71.5** — both beats either alone. ⇒ **vision + text**

PLACEMENT

All layers win

Up / Bottom / All — ImageNet 71.1 / 70.8 / **71.6**; EuroSAT 81.9 / 81.5 / **87.0** (big jump). ⇒ **all layers**

Outcome: our one-seed Figure 3 reproduction supports the paper's default choices — **rank 2 · Q/K/V · both encoders · all layers** — and shows the same trends on ImageNet, StanfordCars, and EuroSAT.

What made reproduction hard

DATA · #1**Dead / incompatible sources**

Caltech101, StanfordCars, EuroSAT and SUN397 had broken or incompatible original sources. We re-sourced the usable datasets via Caltech Data, Kaggle, HuggingFace, or Zenodo.

DATA · #2**SUN397 dropped**

The original source was dead; a HuggingFace version exists, but it is parquet-based rather than the raw JPEG folder structure expected by the original loaders. We therefore report all averages over the 10 reproduced datasets.

SCALE · #3**ImageNet is 144 GB**

Manual login + download of train/val/devkit (~138 GB train alone). The single heaviest item to fetch and extract.

ENV · #4**V100 CUDA/PyTorch mismatch**

Newer PyTorch/CUDA wheels didn't work on V100 out of the box. Fixed with compatible wheels, so we ran on both V100 and A40.

ENV · #5**NumPy ≥ 1.24 break**

`float()` on a 1-D array crashed `cls_acc()`. Fixed by switching to `.item()`.

COMPUTE · #6**Run volume**

150 runs *per* backbone table + ~387 ablation runs (~275 GPU-h total). Solved with a failure-tolerant runner and round-robin job splitting.

Our extension & outlook

Anchoring the adapted model to zero-shot CLIP with a regularization term.



Keep the adapted model close to zero-shot CLIP

Problem. CLIP-LoRA trains with plain **cross-entropy**. With only a few shots, nothing stops the adapters from **drifting away** from CLIP's strong pretrained priors — a likely cause of weaker results on fine-grained sets like **Food101**.

Idea. Add a term that **anchors** the adapted output distribution to the **frozen** zero-shot CLIP — a knowledge-distillation regularizer.

$$L = \text{CE}(\textit{lora_logits}, y) + \lambda \cdot T^2 \cdot \text{KL}(p_{\text{zero-shot}}/T \parallel p_{\text{lora}}/T)$$

- **p_{zero-shot}** — softmax over class similarities from **frozen** CLIP (a fixed teacher).
- **p_{lora}** — softmax from the **adapted** LoRA model; **T** = temperature (**T**² keeps the gradient scale).
- **λ** sets anchor strength; **λ = 0** recovers the CE-only baseline — a clean control. Chosen by ablation: **λ = 0.1, T = 8**.

KL is over class probability distributions — not weight tensors.

What we ran: ablation, then broad evaluation

STEP 1 • ABLATION

Pick λ and T

{Food101, OxfordPets, EuroSAT} $\times$ shots {1,4,16} $\times$ λ {0.1, 0.3, 1.0} $\times$ T {4, 8}, seed 1.

= 54 runs

CHOSEN

$\lambda = 0.1 \cdot T = 8$

The only setting that helps broadly **without hurting the EuroSAT control**; higher λ over-regularizes weak-zero-shot datasets.

STEP 2 • BROAD EVAL

All 10 datasets

Re-ran the full Table-3 grid **with KL on** at $\lambda=0.1$, $T=8$ — 10 datasets $\times$ shots {1,2,4,8,16} $\times$ 3 seeds.

= 150 runs

COMPARE

KL vs CE-only, cell-for-cell

Same grid as Table 3, so every (dataset, shot) is a direct baseline comparison.

CE-only vs KL vs paper – all shots

KL (ours) CE-only paper Δ = KL-CE

ViT-B/16, 10 datasets × 3 seeds, re-run with KL (λ=0.1, T=8). Each cell: KL / CE-only / paper CLIP-LoRA. Top-1 accuracy (%); zero-shot = frozen CLIP. Cells tinted green where KL > CE, red where KL < CE.

Shots	ImgNet	Airc	EuroSAT	Cars	Food	Pets	Flowers	Caltech	DTD	UCF	Avg
0 zero-shot	66.7	23.8	42.1	65.6	85.3	88.2	67.4	93.3	44.5	65.2	64.2
1	70.3 70.4 70.4	29.8 29.8 30.2	73.2 71.1 72.3	70.0 69.8 70.1	85.4 84.7 84.3	91.7 91.9 92.3	81.3 82.3 83.2	94.0 93.9 93.7	54.6 54.0 54.3	75.4 76.0 76.3	72.6 72.4 72.7 +0.19
2	70.9 70.8 70.8	31.9 32.6 33.2	83.3 82.0 82.7	73.4 73.4 73.2	85.2 83.3 83.2	91.2 90.7 91.3	90.0 89.8 89.8	95.0 94.8 94.6	59.6 59.9 59.9	80.1 80.1 80.0	76.1 75.7 75.9 +0.30
4	71.7 71.5 71.4	37.0 38.0 37.9	86.3 85.2 84.9	77.4 77.1 77.4	85.3 83.0 82.7	91.6 90.5 91.0	94.0 93.3 93.7	95.5 95.0 95.2	64.6 63.7 63.8	81.5 81.1 81.1	78.5 77.8 77.9 +0.65
8	72.6 72.4 72.3	44.6 45.4 45.7	90.2 89.1 89.7	82.1 82.1 82.1	85.7 83.1 83.1	92.7 91.3 91.7	96.1 96.1 96.3	96.0 95.7 95.6	68.4 67.5 67.5	84.3 84.0 84.1	81.3 80.7 80.8 +0.59
16	73.8 73.6 73.6	54.7 54.9 54.7	92.8 92.3 92.1	85.9 86.3 86.3	86.3 84.1 84.2	93.4 92.3 92.4	98.1 97.9 98.0	96.5 96.2 96.4	74.0 71.8 72.0	86.8 86.4 86.7	84.2 83.6 83.6 +0.63

Verdict: our CE reproduction matches the paper; KL adds a consistent **+0.2→+0.6** on top (35/50 cells), largest on the fine-grained failure cases — Food101 +2.1, DTD +2.2 at 16-shot — so KL improves on the original CLIP-LoRA too.

Summary – what we did, and what's next

REPRODUCED • HOW

Faithful, end-to-end

Forked the authors' repo + our own runner scripts; Tables 3/4/5 = 10 datasets × 5 shots × 3 seeds × 3 backbones (450 runs), plus a one-seed Figure 3 trend reproduction on ImageNet, StanfordCars, EuroSAT (~275 GPU-h). Every average within **≈0.3%** of the paper across the full grid.

CHALLENGES

Mostly data & scale

Dead dataset URLs, 144 GB ImageNet, CUDA/PyTorch issues, and a NumPy break — handled with dataset fixes, compatible wheels (running on both V100 and A40), and a failure-tolerant runner.

EXTENDED

KL anchor to zero-shot CLIP

$L = CE + \lambda \cdot T^2 \cdot KL(\text{zero-shot} \parallel \text{LoRA})$; ablation picked $\lambda=0.1$, $T=8$. Reduces drift from CLIP's prior.

EXTENSION • RESULT

+0.47 avg · 35/50 cells improved

Biggest gains on Food101 +1.9, EuroSAT +1.2, DTD +0.9 — the failure cases. Supports the hypothesis.

FUTURE WORK

Add SUN397 → full 11-dataset set

3 seeds for Figure 3 (error bars)

Per-dataset λ / T tuning

KL on ViT-B/32 & ViT-L/14

Compare newer PEFT variants: DoRA / LoKi

★ THANK YOU

Questions & discussion

PAPERS

Zanella & Ben Ayed (2024) — *Low-Rank Few-Shot Adaptation of Vision-Language Models*. arXiv:2405.18541

Radford et al. (2021) — *Learning Transferable Visual Models From Natural Language Supervision (CLIP)*. arXiv:2103.00020

CODE

Base: github.com/MaxZanella/CLIP-LoRA

Our repository: github.com/bnkglu/tuw-dl4nlp-2026s

Hasan Berke Bankoglu (12432802) · Zarek Asif (12432949)
e12432802@student.tuwien.ac.at
192.039 — DL4NLP · 2026S

Paper vs our reproduction vs our KL

ViT-B/16, top-1 accuracy (%) averaged over the 10 reproduced datasets × 3 seeds. Paper = Zanella & Ben Ayed (2024), Table 3.

Shots	Paper CLIP-LoRA	Our CE-only	Our KL	Δ KL-CE	Δ KL-paper
1	72.7	72.4	72.6	+0.2	-0.1
2	75.9	75.7	76.1	+0.4	+0.2
4	77.9	77.8	78.5	+0.7	+0.6
8	80.8	80.7	81.3	+0.6	+0.5
16	83.6	83.6	84.2	+0.6	+0.6
avg	78.2	78.0	78.5	+0.50	+0.4

Our CE reproduction matches the paper to ≈0.1 avg; KL adds **+0.47** over our CE and **+0.3** over the paper's original numbers, growing with shots.

Efficient in our hands too

Per-run wall-clock from our logs. Paper reports CLIP-LoRA at **~50 min** for 16-shot ImageNet on a single **A100**; we ran on a mix of **V100 / A40**.

16-SHOT IMAGENET

Paper 50 min (A100)

Ours (V100/A40): **104 min** ViT-B/16 · 92 min ViT-B/32 · 191 min ViT-L/14.

SCALING · IMAGENET VIT-B/16

12 → 104 min

1 / 2 / 4 / 8 / 16-shot ≈ 12 / 19 / 31 / 55 / 104 min — smooth with shots.

WHOLE PROJECT

~275 GPU-h

1041 runs across all tables, Figure 3 and the extension (~11.4 GPU-days).

Same story as the paper — minutes, not hours. Our ~2× on 16-shot ImageNet vs the paper's 50 min is the V100/A40-vs-A100 hardware gap, not a method slowdown; CLIP-LoRA stays far below prompt-tuning methods (e.g. PLOT++ ≈ 15 h on the same task).